

# Reducing Network Overhead in Web Applications

Adam Derewecki  
given for EECS @ CWRU

# \$ whoami

- Graduated BSE Comp Eng 2007
  - Research Assistant for Networking Dept
  - TA for ENGR131 (Fall '05, '06)
- IBM Software Engineer, 2008 - 2009
  - DB2 XML Storage and Runtime teams
- Yelp Backend Engineer, 2009 - Present

# It's only 30KB, will it make a diff?

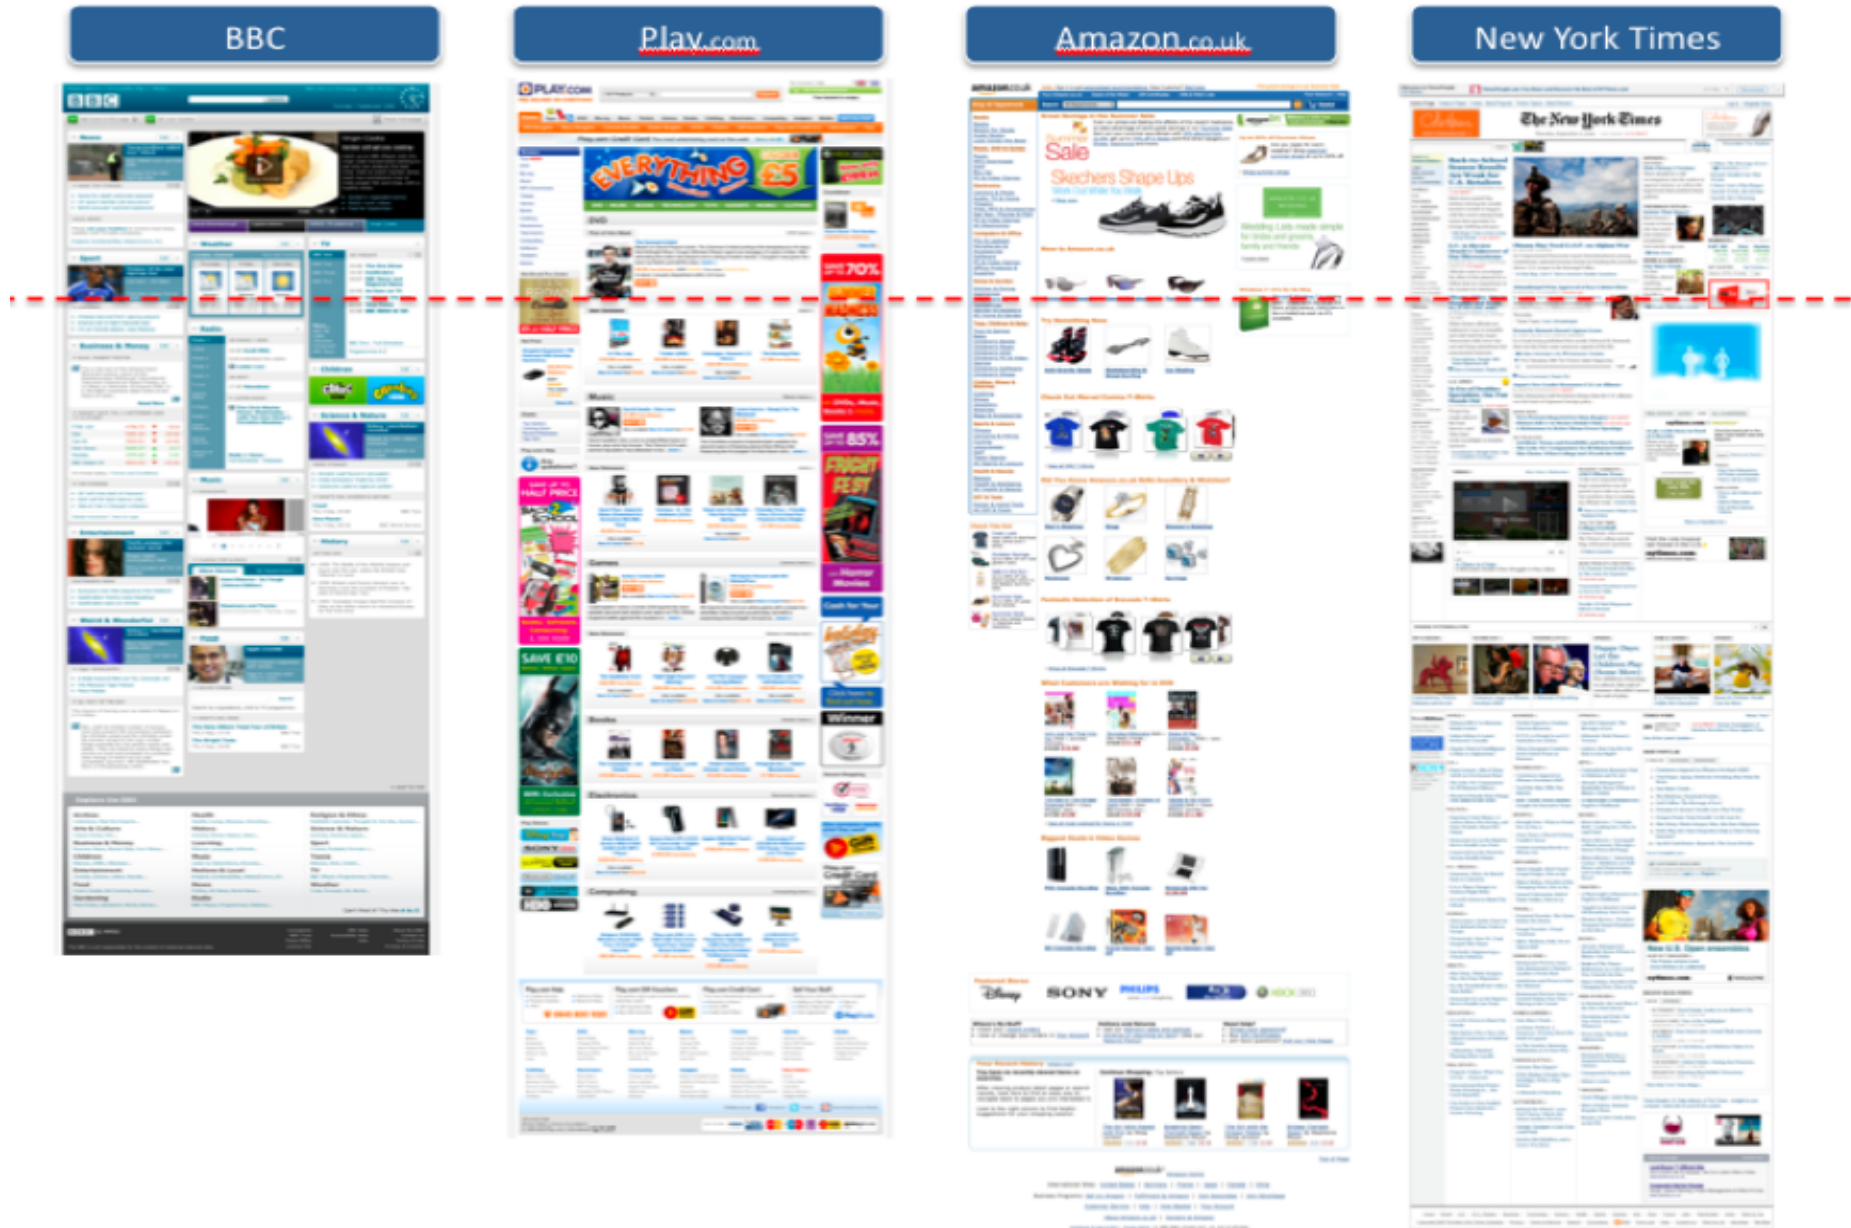
- Hell I'll just cache it!
  - YDN caching experiment\*
    - 40-60% of user have an empty cache
    - 20% of all page views are empty cache
- Packets are 1460 bytes
  - Two packets for 1461 byte file
  - TCP slow start penalty for >2 packets
- Backbones are fast, last mile isn't
  - EDGE: Best case 150ms, real world average 500ms
- In short: yes, it matters.

\* [http://developer.yahoo.com/blogs/ydn/posts/2009/10/a\\_engineers\\_gui/](http://developer.yahoo.com/blogs/ydn/posts/2009/10/a_engineers_gui/)  
<http://yuiblog.com/blog/2007/01/04/performance-research-part-2/>

# AJAX: Don't Repeat Yourself

- Asynchronous JavaScript and XML
  - (doesn't require XML payload)
- Lets browser talk to server via JavaScript after page has finished loading for client
- Why AJAX?
  - load new pages without full page loads
    - search results
    - Google Instant Search is an extreme example
  - don't load more than you have to - lazy loading
    - good for slow components (mint.com)
    - optional components
      - "below the fold"
  - BigPipe
    - <http://www.facebook.com/notes/facebook-engineering/bigpipe-pipelining-web-pages-for-high-performance/389414033919>

# Below the Fold



# Rolling your own XMLHttpRequest

```
// instantiate an "AJAX" object  
var xml = new XMLHttpRequest();
```

```
// set up the connection  
xml.open('GET', '/some/url', true);  
      ^      ^      Async (true) or sync (false)*  
      |      |URL: must be on the same server for security reasons  
      |protocol: GET or POST. Remember, GET has size restrictions
```

```
// send the data payload  
xml.send(data); // slam it out there  
// Note: While it's AJA-XML, the data sent does not need to be XML. // We commonly use  
JSON and HTML snippets.  
// Note: xml.send() is also valid (for example, tracking)
```

Why use synchronous?

- Sometimes a blocking call from a page is useful
- Expire a token or some other final page action before bouncing

# AJAX: Yelp Search Results

Tacos San Francisco

http://www.yelp.com/search?find\_desc=tacos&ns=1&find\_loc=san+francisco,+ca

Now in the UK! Friends' Activity **100** Logged in as Adam D. My Account Log Out

**yelp** Real people. Real reviews.

Search for (e.g. taco, cheap dinner, Max's) **tacos** Near (Address, Neighborhood, City, State or Zip) **san francisco, ca** Search

Welcome About Me Write a Review Find Reviews Invite Friends Messaging Talk Events Member Search

**live On** safely connecting you in ways you never thought possible see more >

**tacos San Francisco** 1 to 10 of 1494 - Results per page: 10 ▾

Hide Filters

Sort By	Neighborhoods	Distance	Features	Price	Category
» Best Match Highest Rated Most Reviewed	<input type="checkbox"/> Mission <input type="checkbox"/> SOMA <input type="checkbox"/> Financial District <input type="checkbox"/> Civic Center/Tenderloin ... More Neighborhoods >	» Bird's-eye View Driving (5 mi.) Biking (2 mi.) Walking (1 mi.) Within 4 blocks	<input type="checkbox"/> Open Now (1:38 am) <input type="checkbox"/> Good for Groups <input type="checkbox"/> Good for Kids <input type="checkbox"/> Take-out ... More features >	<input type="checkbox"/> \$\$\$\$ <input type="checkbox"/> \$\$\$ <input type="checkbox"/> \$\$ <input type="checkbox"/> \$	<input type="checkbox"/> Mexican <input type="checkbox"/> Food Stands <input type="checkbox"/> Fast Food <input type="checkbox"/> Bars ... More categories >

1. **Tacos El Primo**  
Categories: Mexican, Food Stands  
Neighborhood: Bayview/Hunters Point

★★★★★ 18 reviews  
Reviewed by: 1 friend  
Yosemite & Jennings  
San Francisco, CA 94124  
(415) 846-4975

Hot damn, son... this just might be the best **taco** shop I've ever eaten at; and not just in San Francisco, but anywhere else in the world (including Taiwan!). Growing up in San Diego, I was blessed

2. **San Buena Taco Truck**  
Categories: Mexican, Food Stands  
Neighborhood: Mission

★★★★★ 262 reviews  
Reviewed by: 1 friend  
22nd St & Harrison St  
San Francisco, CA 94110  
(415) 559-6127

Visiting the Mission District with my brother, this was one of our favorite stops. The **tacos** were amazing, inspiring me to try to find a recipe for **tacos** al pastor to match. Highly recommended

3. **San Buena Taco Truck**  
Categories: Food Stands, Mexican  
Neighborhood: Mission

★★★★★ 22 reviews  
Shotwell and 16th St  
San Francisco, CA 94199  
(415) 559-6127

Another favorite spot of mine for probably the past decade now. I do the **tacos**; carne asada or pork. Always get lime on the side too. Cash only, but only \$1.75/**taco**. They do to go too so you don't

4. **Don Pisto's**  
Categories: Mexican, Breakfast & Brunch  
Neighborhood: North Beach/Telegraph Hill

★★★★★ 175 reviews  
510 Union St  
San Francisco, CA 94133  
(415) 395-0939

Upcoming Event

At this point I have ordered almost everything on the menu and I have loved every bite. The carne asada **tacos** are my favorite and the fish **tacos** are always perfection. The tortilla soup is the best

Mo' Map  
☐ Map, stay put! ☐ Redo search in map

POWERED BY Google Maps data © 2010 Google - Terms of Use

derwiki... Buddy List sysinfra... [Yelp Mai... Reducin... Downloa... Tacos S... [Develo... derwiki

1:38

# AJAX: Yelp's 'the hub'

- AJAX heavy web pages have JS that get messy
  - tangled webs of chained callbacks
- 'the hub'
  - mediocre name, useful technology
  - Create a model of all the interactive elements on the page
  - Publish messages to the hub when things finish
  - Subscribe to do things when messages are published
  - sure beats daisy chains of callbacks
  - <http://github.com/Yelp/thehub/blob/master/js/thehub.js>



# AJAX: Yelp's PubSubHub

```
deals.queueDeal = function(event) {
  deals.doFunctionOnSaveSuccess('admin.deals.queue_status', deals.QueueStatus, function() {
    yelp.hub.publish('admin.deals.queue_status', new deals.QueueStatus('IN_PROGRESS'));
    deals.queueRequest = new yelp.async.ajax.Request('/deal/queue_deal_to_send/' + deals.
getDealId(), {
  method: 'post',
  onSuccess: function(resp) {
    var respJSON = resp.responseJSON;
    if (resp.responseText) {
      console.log(resp.responseText);
    }
    if (respJSON.success) {
      yelp.hub.publish('admin.deals.queue_status', new deals.QueueStatus
('COMPLETE_SUCCESS'));
    } else {
      var error_msgs = deals.collectErrorMessages(respJSON);
      yelp.hub.publish('admin.deals.queue_status', new deals.QueueStatus('COMPLETE_FAIL',
error_msgs));
    }
  },
  onFailure: function(resp) {
    yelp.hub.publish('admin.deals.save_status', new deals.QueueStatus('COMPLETE_FAIL', 'Show
this to an engineer - code ' + resp.status));
  }
});
});
};
```

# AJAX: Yelp's PubSubHub


```
deals.QueueStatus = function(state, errorMessage) {  
  var states = {'IN_PROGRESS': {'success': null, 'message': 'Queuing...'},  
    'COMPLETE_FAIL': {'success': false, 'message': 'Failed to add to queue: ' + errorMessage},  
    'COMPLETE_TIMEOUT': {'success': false, 'message': 'Failed to add to queue: save timed out.'},  
    'COMPLETE_SUCCESS': {'success': true, 'message': 'Deal queued!'}};  
  this.type = 'QUEUE_STATUS';  
  yelp.pages.admin.deals.initStatus(this, states, state);  
};
```

```
deals.observeButtons = function() {  
  $(document.body).delegate('click', '#queue_deal', deals.queueDeal)  
  .delegate('click', '#dequeue_deal', deals.dequeueDeal)  
};
```

```
deals.hubSubscribe = function() {  
  yelp.hub.subscribe('admin.deals.queue_status', deals.updatePageStatus);  
};
```

# One more real quick AJAX thing

- Staged Loading
  - Load core JS in one file
  - Have JS load additional, non-essential JS after the main page has loaded
- Inline JS
  - JavaScript is single threaded, blocking
  - Hiding JS in a comment and eval'ing it
  - <http://www.stevesouders.com/blog/2009/12/07/downloading-javascript-as-strings/>



**i m in ur internet**

**cloging ur tubes**

# Longpolling

- AJAX is cool and all -- but how do you push dynamic content to the client?
- You can't -- but you can!
  1. Have the browser open a long running AJAX request to a webserver
  2. The web server holds the connection until it has something to say back to the client
  3. After the server pushes a notification back to the client, the client starts a new long running AJAX request to the server to wait for more updates
- This is how push services (or at least Yelp's) are done
  - "I'm still here" every 15 minutes
  - High water mark, 100,000 active connections

# Longpolling: Server considerations

- Traditional polling methods don't keep open connections, server doesn't need to keep as many at once
- Per-process/per-thread models tend to have a lot of extra weight for each additional connection
- Solution: epoll!
  - asynchronous sockets (non-blocking)
  - single event loop instead of process/thread models
  - each iteration of the loop looks at a set of file descriptors (sockets) you have interest in, and handles a bit of reading or writing on each of them
  - Tornado and node.js are two popular frameworks that let you handle I/O this way

# epoll server

```
import socket, select
serversocket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
serversocket.setsockopt(socket.SOL_SOCKET, socket.SO_REUSEADDR, 1)
serversocket.bind(('0.0.0.0', 9090))
serversocket.listen(1)
serversocket.setblocking(0) # nonblocking
epoll = select.epoll()
epoll.register(serversocket.fileno(), select.EPOLLIN)
default_resp = str(range(1000))
cxns, reqs, resps = {}, {}, {}
try:
    while True:
        events = epoll.poll(1)
        for fileno, ev in events:
            if fileno == serversocket.fileno():
                # new cxn
                conn, addr = serversocket.accept()
                conn.setblocking(0)
                epoll.register(conn.fileno(), select.EPOLLIN)
                cxns[conn.fileno()] = conn
                reqs[conn.fileno()] = ""
                resps[conn.fileno()] = default_resp
            elif ev & select.EPOLLIN:
                reqs[fileno] += cxns[fileno].recv(1024)
                if any(eol in reqs[fileno] for eol in (b'\n\n', b'\n\r\n')):
                    epoll.modify(fileno, select.EPOLLOUT)
                    print reqs[fileno].decode()[:-2]
            elif ev & select.EPOLLOUT:
                bytes_written = cxns[fileno].send(resps[fileno])
                resps[fileno] = resps[fileno][bytes_written:]
                if len(resps[fileno]) == 0:
                    epoll.modify(fileno, 0)
                    cxns[fileno].shutdown(socket.SHUT_RDWR)
            elif ev & select.EPOLLHUP:
                epoll.unregister(fileno)
                cxns[fileno].close()
                del cxns[fileno]
finally:
    epoll.unregister(serversocket.fileno())
    serversocket.close()
```

**Teh Intertubes**




**They is clogged**




# epoll resources

- <http://scotdoyle.com/python-epoll-howto.html>
  - previous slide borrowed heavily from scot's code
- <http://sheddingbikes.com/posts/1280829388.html>
  - Zed Shaw, of internet fame
- Frameworks
  - <http://github.com/facebook/tornado> (Python)
    - developed at FriendFeed (including two CWRU grads!)
  - <http://nodejs.org/> (server-side JS)
  - <http://github.com/eventmachine/eventmachine> (Ruby)

# nginx (epoll) vs Apache: speed



# nginx vs. Apache: memory

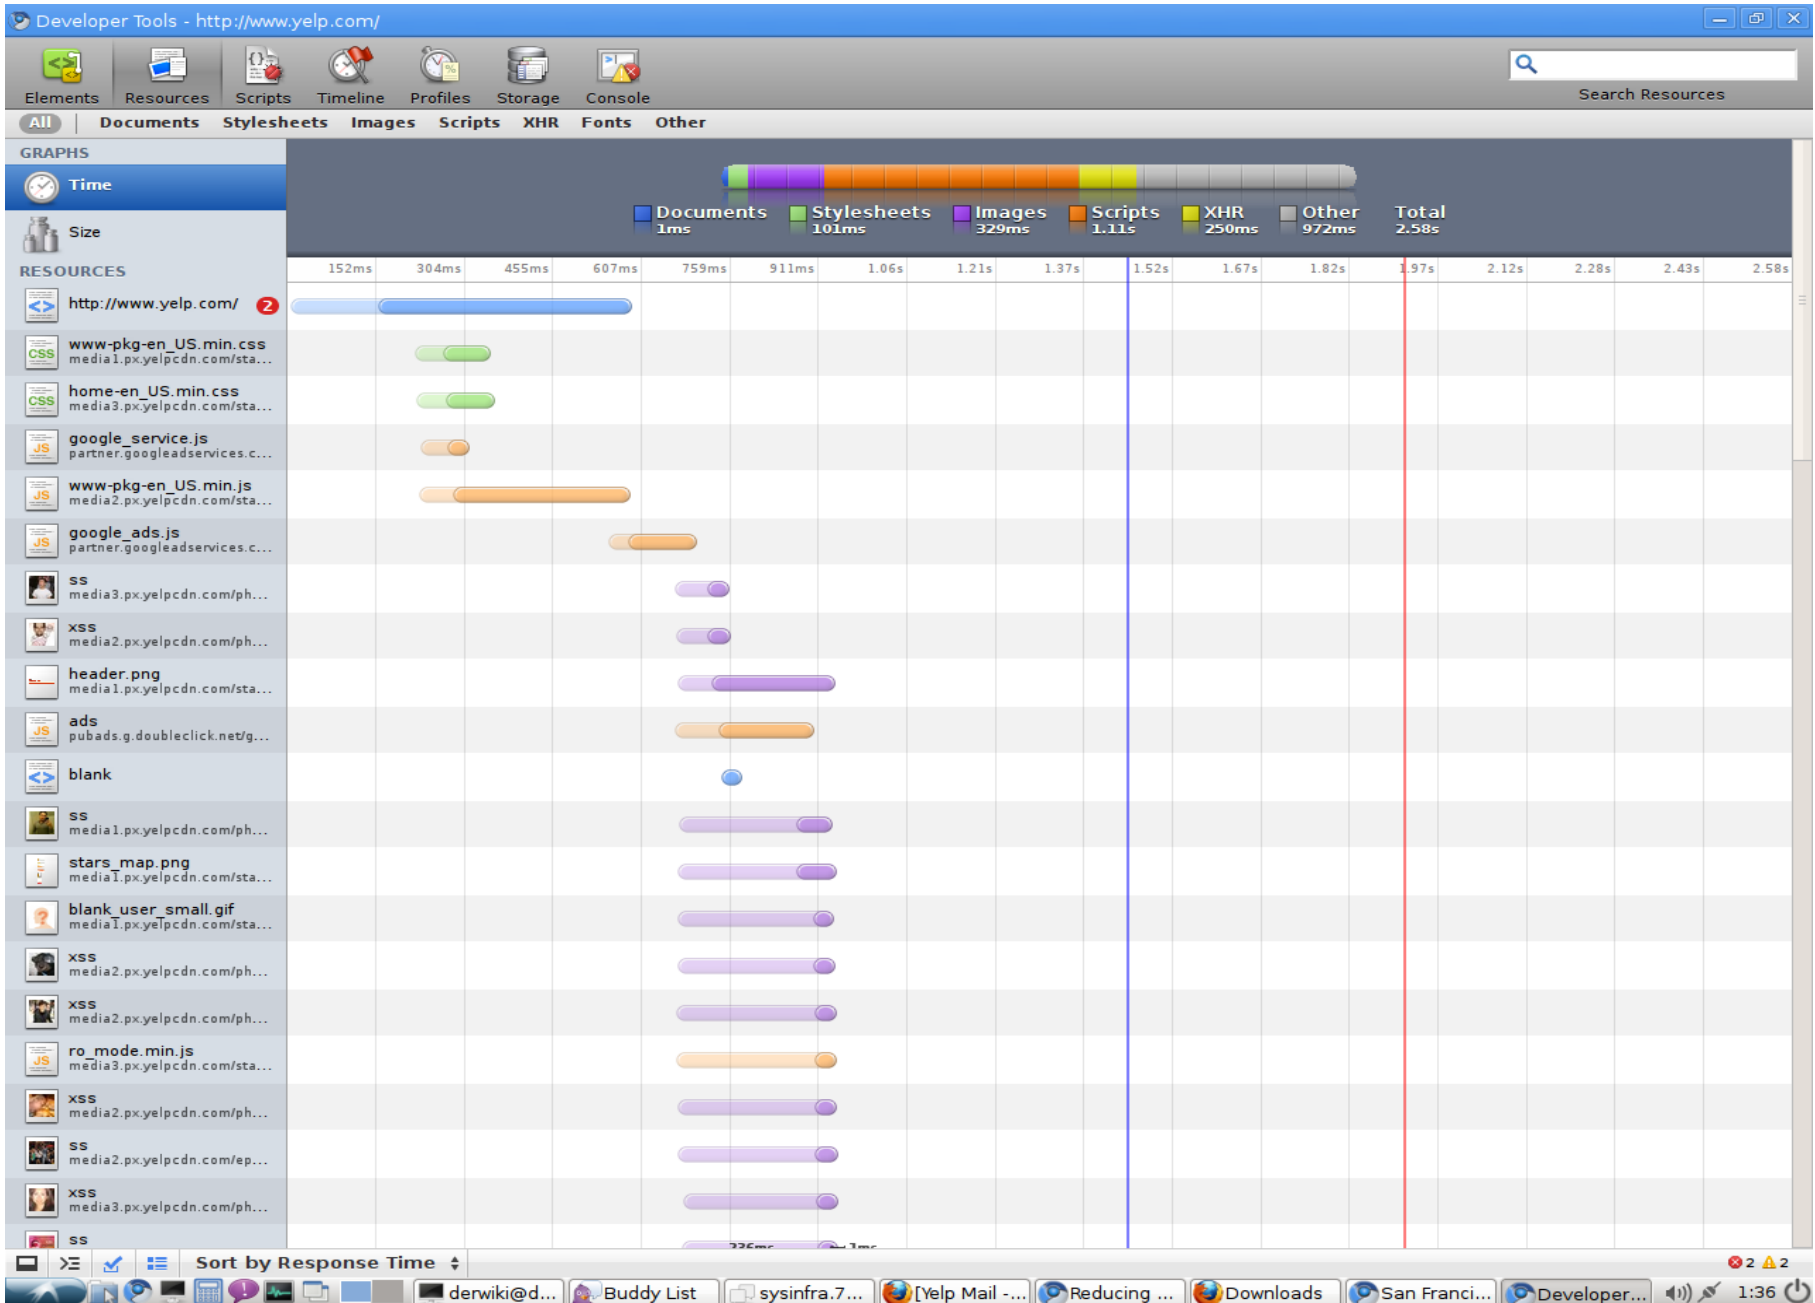


- Ouch, that hurts!

# Longpolling: easier on the network

- 100,000 concurrent connections
- Average idle time of 5 minutes
- Poll every 30 seconds
  - 3333 requests per second
- Longpoll, only take action every 5 minutes
  - 333 requests per second
- Longpolling combined with epoll reduces network traffic and also scaled better than traditional poll-thread/process systems

# Measuring front end performance



# Making your page load faster

- Gzip. Gzip. Gzip.
  - <http://paulbuchheit.blogspot.com/2009/04/make-your-site-faster-and-cheaper-to.html>
- Reduce number of assets you serve
- ... and the size of them
  - spriting images
  - combining JS/CSS
  - minifying JS/CSS
- Reduce number of DNS lookups
- JavaScript loading is a blocking operation
  - only load essential JS when you need to
- Performance tools
  - Chrome, Safari have built in tools
  - Firebug ([www.getfirebug.com](http://www.getfirebug.com))
  - Yahoo YSlow ([developer.yahoo.com/yslow](http://developer.yahoo.com/yslow))

# Cookies

- If you give an [asset] a cookie...
  - He'll send it along every request to your domain
  - Like real cookies, those calories (bytes) add up
  - Use another domain name for static content
    - [www.yelpcdn.com](http://www.yelpcdn.com)
- Test sites yourself
- `javascript:alert(document.cookie.length);`
  - [www.yelp.com](http://www.yelp.com) - 1985 bytes
  - [www.facebook.com](http://www.facebook.com) - 450 bytes
  - [www.twitter.com](http://www.twitter.com) - 1123 bytes
  - [www.google.com](http://www.google.com) - 401 bytes



# Spriting



- Combine 20 images used frequently into one
- Reduces number of HTTP requests
- Caching (during this session at least)
- CSS to crop to what you need

```
img.stars_3 {  
    clip: rect(95px 83px 110px 0px);  
    top: -95px;  
}
```



# The Final Slide

- Any of this sound interesting? We're hiring!
  - [www.yelp.com/careers](http://www.yelp.com/careers)
  - 50 person engineering team
  - Open source
    - [github.com/yelp](https://github.com/yelp)
    - [engineeringblog.yelp.com](http://engineeringblog.yelp.com)
  - Located in downtown San Francisco

